

Reconstruction of objects from grasps with latent diffusion modeling

Christian Yankacar

Nikola Jankulovic

Abstract—The goal of this project is to design and train a latent diffusion model that takes geometric information from a anthropomorphic robotics hand grasping an object as an conditional input and reconstructs said object based on the information provided by the input. Research questions of this project are, how much geometric information is encoded in the grasp of a robotic hand and how can this information be used to reconstruct the grasped object. Which latent space encoding and variational autoencoder architecture achieves the best compromise between preserving the information of the original input data and accelerating the training of the model. Which diffusion model architecture offers the best reconstruction quality. Which conditioning technique is best suited for guiding the diffusion process to generate object shapes based on the grasp input. How should the training data be preprocessed or enhanced to improve the training performance of the latent diffusion model.

Keywords—*Latent Space, Diffusion, 3d-object reconstruction, Generative modeling, Humanoid robotics*

I. INTRODUCTION

Grasping objects is one of the most important tasks in robotics. Recent data-driven approaches tried to solve this task by synthesizing grasps with generative artificial intelligence (AI) like variational autoencoders (VAE) or diffusion models from the geometry of the object as an input [1–3]. While being effective in generating a diverse set of possible grasps, these models also generate many invalid grasps and often rely on optimization techniques or discriminator networks to detect invalid grasps. This problem arises from the incomplete geometric representation of the input, because only the front of the object is seen by the robot. To improve the success rate of valid grasp generation recent works investigated neural shape completion techniques from partial geometric input [4, 5]. These works relied on vision-based systems to provide the partial geometric information. But for industrial applications robots have to be cheap and are often not equipped with camera systems. Furthermore, robots in an industrial setting often have to work under bad lighting conditions or in small occluded work spaces, where vision cannot be provided. To resolve this issue other works explored the idea of reconstructing object shapes from tactile sensing [6, 7]. These works use tactile sensors that generate a pressure image of the contact surface from which they extract features with transformer or convolution networks. We are interested in the task of shape completion based from the joint angles of the robot hand, when it is grasping the object. So, we focus on the geometric information that is inherently coded into the hand-object interaction without the reliance on tactile sensors. The closest work that investigated this idea is UGG from [8]. In UGG the authors introduced a unified latent space where object O and grasp G where modeled as a joint probability $P(G, O)$. The authors showed that this latent space allows for training both the conditional generation of objects and grasps.

While this approach is interesting, we want to develop a direct strategy by directly training the diffusion model to learn the conditional probability of generating objects from a given grasp $P(O|G)$ and compare our results with UGG.

II. TECHNICAL OUTLINE

Our main idea is inspired by works like [9, 10] who showed that the DDPM [11] process can be performed in the latent space to accelerate the training, while retaining state of the art generation quality. As we are dealing with high-dimensional 3d data and have limited computing resources latent diffusion will be key in giving us good generation quality for reasonable computational costs. In order to further reduce the computational costs, we decided to represent the objects as point clouds. Recent works have proven that the training and reconstruction of point clouds with diffusion models achieves very good results [8, 12, 13]. For latent diffusion the VAE is of great importance as it generates the latent code to perform the diffusion on. For our work we would like to experiment with different VAEs. Pointnet++ and PVCNN [14, 15] are popular and easy to use autoencoders that generate global shape latents. LION [12] a more advanced architecture constructs both a global shape latent and a point structure latent that better retains the geometrical structure of the original point cloud, thus providing the diffusion model with more geometric information. Then the diffusion model itself must be considered. Traditional latent diffusion models leverage a U-Net Design that relies on convolution layers [9, 10], but there was also much work done on diffusion transformers for point cloud generation recently [7, 13, 16]. These works showed that transformers are very good in preserving the shape of the original object by tokenizing the object point cloud and using positional encoding on the tokens. We will also investigate different input modalities besides the joint angles of the robot hand like for instance a point cloud representation of the robot's hand and its effect on reconstruction quality and training time. Different conditioning methods like simple concatenation, FiLM or cross-attention will also be evaluated [17, 18]. But for the start of the project, we will keep the model as simple as possible to have a solid baseline that can be improved upon. For the first two weeks we will explore the provided training data and decide which augmentations should be used if needed. Then we will start with a simple latent diffusion architecture that is unconditionally trained to reconstruct point clouds. For this we will rely on pretrained VAE first and test if this yields good results. This simple testing in the early stages also serves to compare U-Net based and transformer-based diffusion models. Based on the test results we are going to choose a model architecture and iteratively add more features like grasp conditioning to develop our model.

III. REFERENCES

- [1] H.-S. Fang, C. Wang, M. Gou, and C. Lu, "GraspNet: A Large-Scale Clustered and Densely Annotated Dataset for Object Grasping," 2019. [Online]. Available: <https://arxiv.org/pdf/1912.13470>
- [2] J. Carvalho *et al.*, "Grasp Diffusion Network: Learning Grasp Generators from Partial Point Clouds with Diffusion Models in $SO(3) \times R^3$," 2024. [Online]. Available: <https://arxiv.org/pdf/2412.08398>
- [3] K. R. Barad, A. Orsula, A. Richard, J. Dentler, M. A. Olivares-Mendez, and C. Martinez, "GraspLDM: Generative 6-DoF Grasp Synthesis Using Latent Diffusion Models," *IEEE Access*, vol. 12, pp. 164621–164633, 2024, doi: 10.1109/ACCESS.2024.3492118.
- [4] D. Hidalgo-Carvajal *et al.*, "Anthropomorphic Grasping With Neural Object Shape Completion," *IEEE Robot. Autom. Lett.*, vol. 8, no. 12, pp. 8034–8041, 2023, doi: 10.1109/LRA.2023.3322086.
- [5] S. Iwase *et al.*, "ZeroGrasp: Zero-Shot Shape Reconstruction Enabled Robotic Grasping," 2025. [Online]. Available: <https://arxiv.org/pdf/2504.10857>
- [6] L. Gu, H.-J. Huang, M. Qadri, M. Kaess, and W. Yuan, "TouchAnything: Diffusion-Guided 3D Reconstruction from Sparse Robot Touches," 2026. [Online]. Available: <https://arxiv.org/pdf/2604.08945>
- [7] H. Zhang, X. Zhang, J. Huang, Z. Feng, and X. Xiao, "End-to-End Diffusion-Based 3D Object Reconstruction From Robotic Tactile Sensing," *IEEE Robot. Autom. Lett.*, vol. 11, no. 2, pp. 1434–1441, 2026, doi: 10.1109/LRA.2025.3641115.
- [8] "UGG Unified Generative Grasping,"
- [9] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, "High-Resolution Image Synthesis with Latent Diffusion Models," 2021. [Online]. Available: <https://arxiv.org/pdf/2112.10752>
- [10] D. Podell *et al.*, "SDXL: Improving Latent Diffusion Models for High-Resolution Image Synthesis," 2023. [Online]. Available: <https://arxiv.org/pdf/2307.01952>
- [11] "Denoising Diffusion Probabilistic Models,"
- [12] "LION Latent Point Diffusion Models for 3D Shape Generation,"
- [13] J. Ji, R. Zhao, and M. Lei, "Latent diffusion transformer for point cloud generation," *Vis Comput*, vol. 40, no. 6, pp. 3903–3917, 2024, doi: 10.1007/s00371-024-03396-1.
- [14] "PointNetpp,"
- [15] Z. Liu, H. Tang, Y. Lin, and S. Han, "Point-Voxel CNN for Efficient 3D Deep Learning," 2019. [Online]. Available: <https://arxiv.org/pdf/1907.03739>
- [16] "DiT-3D Exploring Plain Diffusion Transformers for 3D Shape Generation,"
- [17] "Attention is all you need,"
- [18] Ethan Perez, Florian Strub, Harm de Vries, Vincent Dumoulin, and Aaron Courville, "FiLM: Visual Reasoning with a General Conditioning Layer,"